

SOLUTION ARCHITECTURE · POC PROPOSAL

DPDP Compliance Platform on Databricks

*A unified, evidence-producing platform for
operating under the Digital Personal Data
Protection Act — six modules, six workflows, one
audit trail.*

Prepared by · Databricks solution architecture

For · Chief Compliance Officer, General Counsel, Chief Financial Officer

Status · Working draft — pre-POC scoping

Date · 17 April 2026

Executive summary

The Digital Personal Data Protection Act exposes every Indian enterprise to penalties of up to ₹250 crore per contravention, a 72-hour breach notification deadline, and data principal rights that must be fulfilled against every record the organization holds. The operational reality is that most enterprises today cannot locate their own personal data with confidence, cannot evidence consent at the granularity the Act requires, and cannot reproduce historical decisions under regulator inspection.

This document proposes a six-module platform on Databricks that addresses all six core DPDP obligations — PII discovery, consent intelligence, data subject rights, breach detection, compliance audit, and retention with cross-border monitoring — under a single Unity Catalog governance spine. The design incorporates three recent Databricks capabilities that meaningfully improve the architecture over a 2025-era approach: **Lakewatch** (agentic SIEM) collapses breach response from 72 hours to under 18; **Lakebase** (Postgres OLTP) provides the transactional tier that consent capture and DSR intake genuinely need; **Agent Bricks** automates the document-generation work — DPIAs, DPBI notifications, multi-language notices, inspection packets — that otherwise consumes weeks of Legal time.

The proposal closes with a 14-day rapid discovery sprint against a single source system, producing three demonstrable artifacts — a living personal data register, an immutable consent log, and a fulfilled synthetic DSR — that give the Chief Compliance Officer, General Counsel, and Chief Financial Officer enough evidence to approve Phase 1.

Contents

PART I · CONTEXT

1. The DPDP regulatory context
2. The seven-stakeholder map
3. The six-module platform in overview

PART II · ARCHITECTURE

4. Platform architecture at a glance
5. Evolution with Lakewatch, Lakebase and Agent Bricks
6. Module 01 — PII discovery and data inventory
7. Module 02 — Consent intelligence engine
8. Module 03 — Data subject rights hub
9. Module 04 — Breach detection and 72-hour alerting
10. Module 05 — Compliance audit and DPIA
11. Module 06 — Retention, masking and transfers

PART III · REALITIES

12. Module 0 — the organizational prerequisites
13. Friction audit across the six modules
14. Shadow data and the three zones

PART IV · OPERATIONAL WALKTHROUGHS

15. Breach response timeline
16. Compliance audit cycle
17. Single DSR — Priya's walkthrough
18. DPBI Section 28 inspection
19. Retrospective consent campaign
20. New-principal consent — Meera's walkthrough

PART V · COVERAGE

21. Persona journeys across the six workflows
22. Workflow-to-module coverage matrix

PART VI · PROPOSAL

23. POC scope card — 14-day sprint
24. POC archetype comparison
25. Databricks use case classification
26. Phase 1 outline and next steps

PART I

The opportunity

Before designing the platform, we establish what it exists to do. DPDP is not a technology problem; it is an operational obligation that touches every data principal the organization holds, every processing activity it runs, and every stakeholder who signs off on how data is used. This part of the document frames that reality before any architectural decisions are introduced.

1 The DPDP regulatory context

The Digital Personal Data Protection Act of 2023 is the most consequential privacy legislation India has passed, and the first that places enforceable obligations on every organization processing personal data regardless of sector.

The Act's mechanics can be stated briefly. Every organization that processes personal data of Indian residents is a *data fiduciary* with defined obligations: to obtain specific consent for each processing purpose, to honour data principal rights for access, correction, and erasure, to notify the Data Protection Board of India (DPBI) of personal data breaches within 72 hours, to implement reasonable security safeguards, and to conduct annual Data Protection Impact Assessments where designated a Significant Data Fiduciary. Penalties scale up to ₹250 crore per contravention. Significant Data Fiduciary designation, when it arrives via gazette notification, adds the DPIA mandate, independent audit, and the requirement to name a Data Protection Officer.

What makes the Act operationally hard is less its letter and more its scope. The definition of personal data is broad — any information from which a person can be identified directly or indirectly. The consent standard is strict — *free, specific, informed, unconditional*. The rights must be fulfillable against every record the organization holds, including records in systems the compliance team may not even know about. The breach deadline runs against the clock whether or not the organization has yet determined which principals are

affected. The DPIA must reflect current processing activities, not a documentary snapshot from six months earlier.

The gap between what the Act requires and what most Indian enterprises can actually evidence today is substantial. Personal data sits across dozens of source systems; consent is often captured as a T&Cs; acceptance that predates the Act's specificity requirement; data subject rights are handled manually through legal inboxes with multi-week SLAs; breach detection depends on security tools that weren't designed to produce regulator-grade evidence. Closing that gap is the programme this document describes.

2 The seven-stakeholder map

DPDP compliance is not owned by a single function. Seven executive stakeholders have material interests in how the programme runs, and each brings a different success criterion to the table. The platform must serve all seven — not because every stakeholder is equally engaged in every workflow, but because the programme's defensibility depends on each owning their slice.

The Chief Compliance Officer owns the programme as a whole. They track the compliance score, report to the board quarterly, and sign off on annual DPIAs. They are the one persona who is primary or active in every one of the six operational workflows this document covers.

The General Counsel owns legal defensibility. They approve notice templates, adjudicate erasure-versus-retention conflicts, sign off on DPBI submissions. When a Section 28 inspection arrives, they are the primary owner across all four phases of the response — unique to this workflow.

The Chief Information Security Officer owns breach detection and response. Under the new architecture, their toolkit gains Lakewatch's agentic SIEM — this is the persona whose 72-hour workflow compresses most visibly with the platform in place.

The Chief Marketing Officer owns the retrospective consent campaign and the steady-state new-principal consent flow. Their framing is that compliance infrastructure is what makes growth marketing sustainable under DPDP — every new principal acquired through a clean consent flow is a marketing asset that doesn't need cleanup later.

The Chief HR Officer owns employee data — which under DPDP is every bit as sensitive as customer data. Biometric attendance, health insurance enrollments, performance analytics all fall under the Act. The CHO inherits the consent pattern from the customer-facing CMO motion and runs it for the workforce.

The Chief Financial Officer owns the economics. Their framing is penalty avoidance plus the ROI of compliance automation versus legal and audit spend. They are deeply engaged at annual DPIA time and acutely engaged during any DPBI inspection, because that is when the reserves conversation becomes specific.

The Chief Technology Officer owns platform integrity. The architectural discipline that makes the consent schema consistent across every acquisition channel, the retention clocks accurate, the audit log immutable — these are CTO-level invariants that compliance assumes but engineering must enforce.

3 The six-module platform in overview

The platform decomposes DPDP compliance into six modules, each aligned to a specific obligation. The decomposition matters because each module has a distinct owner, cadence, and operational rhythm — and because the Act itself maps cleanly to these six areas.

Module 01 — PII discovery and data inventory builds and maintains the living register of every personal data column, document, and field across the estate. Structured data is classified by Presidio on Spark; unstructured content is classified by Databricks' `ai_classify` and `ai_extract` functions. Unity Catalog holds the tags, lineage, and sensitivity tiers. The output is the personal data register the Act expects every data fiduciary to maintain.

Module 02 — Consent intelligence engine captures every consent event with notice version, language, channel, timestamp, and purpose-level granularity. Writes go to Lakebase for transactional integrity; the immutable audit log lives in Delta; withdrawals propagate through DLT streaming to every downstream system that uses the data.

Module 03 — Data subject rights hub handles access, correction, and erasure requests end-to-end. Unity Catalog lineage discovers where each principal's data lives; Delta DELETE plus VACUUM executes structured erasure; Graph API handles Zone 1 unstructured content with a split execution path for masking, de-indexing, and scheduled residual purge.

Module 04 — Breach detection and 72-hour alerting runs on Lakewatch's agentic SIEM. Detection, affected-principal lookup via lineage, and Agent Bricks-composed DPBI notification drafting compress the full response to under 18 hours — leaving 54 hours of margin inside the Act's deadline.

Module 05 — Compliance audit and DPIA reads signal from every other module, produces a live compliance score, and uses Agent Bricks to generate the annual DPIA automatically from Unity Catalog metadata. A three-month DPIA exercise becomes a three-day regeneration.

Module 06 — Retention, masking and transfers enforces retention policies, manages column-level masking for governed analytics, and monitors cross-border transfers via Unity Catalog lineage. It receives triggers from consent withdrawal, retention expiry, employee exit, and deferred DSR residual purges.

PART II

The architecture

The platform runs on a single governance spine — Unity Catalog — under which six modules consume from and write to a shared lakehouse plus Lakebase OLTP tier. This part walks through the whole-platform view first, then zooms into each module's distinctive mechanism.

4 Platform architecture at a glance

A first-pass layered architecture showing how the seven stakeholders, four consumer interfaces, six DPDP modules, and the governance spine sit above the lakehouse and its upstream sources.

The architecture in figure 4.1 reads top-down. Stakeholders interact through four consumer surfaces — dashboards, the DSR portal, the breach notifier, and the marketing API. These surfaces query the six DPDP modules, which in turn read from and write to the Delta medallion (Bronze/Silver/Gold) under Unity Catalog governance. Processing and ML capabilities (DLT, Structured Streaming, MLflow, Workflows) serve the modules' computational needs. Source systems feed the lakehouse through Auto Loader and JDBC at the base of the diagram.




FIGURE 4.1 Original layered architecture with seven stakeholders, four consumer interfaces, six modules, Unity Catalog governance, and the medallion lakehouse.

5 Evolution with Lakewatch, Lakebase and Agent Bricks

Three Databricks capabilities released in the first quarter of 2026 materially improve the architecture. Figure 5.1 shows the updated platform with each slotted into its natural home.

The updates are targeted, not pervasive. **Lakewatch** replaces the Module 04 detection stack because its Claude-powered agentic SIEM, launched March 2026, correlates structured and unstructured signals more effectively than the MLflow Isolation Forest approach it supersedes. **Lakebase** — Databricks' managed Postgres OLTP layer, generally available as of February 2026 — adds a transactional tier for consent events and DSR intake that the original design handled less cleanly through streaming writes to Delta. **Agent Bricks** becomes the unified agent control plane for all document-generation tasks the platform performs: DPIA drafts, DPBI notifications, multi-language consent notices, and DSR response bundles.




FIGURE 5.1 Updated platform architecture with Lakewatch, Lakebase, Agent Bricks, AI/BI Genie, and ai_classify/ai_extract integrated. The agent orchestration layer is new; Lakebase joins the data tier alongside the medallion.

What does not change is equally important. Unity Catalog remains the single governance spine; the medallion structure continues to serve all analytical workloads; the module decomposition is unaffected. The additions are additive capabilities, not architectural restructuring — which is what makes them deployable against a live compliance programme without a platform rewrite.

6 Module 01 — PII discovery and data inventory

The distinctive mechanism is a dual classification path — Presidio for structured, ai_classify and ai_extract for unstructured — converging on Unity Catalog's tag and lineage layer. The ingestion layer evolves across three versions; we present the final v3 design below.

Module 01 builds and maintains the personal data register that every downstream module depends on. The challenge it solves is organizational as much as technical: most Indian enterprises have 40+ source systems holding personal data, each owned by a different team, and coordinating credential distribution across them is where classification projects stall before they start.

The final design handles this by offering **five ingestion patterns** rather than one. Lakehouse Federation for databases (Postgres, Snowflake, BigQuery) queries in place without data movement. Lakeflow Connect for SaaS (Salesforce, Workday, ServiceNow) uses managed connectors. A landing-zone pattern accepts CSV snapshots dropped into a named S3 prefix by source teams, sidestepping credential sharing entirely. Delta Sharing handles Databricks-native upstreams. A Zone 1 tier — SharePoint, OneDrive, Exchange, Google Workspace, Teams archives — covers the unstructured enterprise collaboration surface where most real-world PII hides. An exclusion policy gate sits between ingestion and Bronze, filtering out privileged content, HR-sensitive material, and legal-hold records before classification runs.




FIGURE 6.1 Module 01 v3 with five ingestion patterns, exclusion policy gate, dual classification (Presidio + ai_classify/ai_extract), external classifier inputs, and the living personal data register as output.

External classifier inputs — Microsoft Purview, Varonis, BigID — feed into Silver alongside the two native classifiers, so customers who have invested in these tools get their existing classifications respected rather than overwritten. Agent Bricks resolves conflicts when the three classifiers disagree.

7 Module 02 — Consent intelligence engine

The distinctive mechanism is real-time withdrawal propagation. One click from a customer must fan out to every downstream system that processes their data, in minutes rather than the daily batch cadence most marketing stacks use today.

Consent events arrive from four channels — web, mobile app, call center, partner API. All four write to a single Lakebase endpoint with an identical schema; this is the non-negotiable design invariant that prevents the six-schema drift common in multi-channel consent systems. The Lakebase OLTP tier gives the platform read-after-write consistency and atomic commits — which matters because every consent event is actually three or more discrete decisions (marketing, analytics, third-party sharing), committed as one transaction.




FIGURE 7.1 Module 02 consent engine: four capture channels write to Lakebase; Delta sync creates the immutable log; DLT CDC stream propagates withdrawals in real time to downstream systems (marketing CRM, analytics, feature store, partners).

The Delta immutable log is the evidentiary layer. Every event is stamped with notice version, language, channel, timestamp, IP, and device. Three years later, when a data principal disputes the validity of a consent, the GC can reproduce the exact notice they saw on the exact date in their exact language. DLT CDC then ensures withdrawals propagate through downstream Gold tables in under five minutes, so a 10pm withdrawal is out of every active campaign before morning.

8 Module 03 — Data subject rights hub

The distinctive mechanism is Unity Catalog lineage-based discovery. One query locates every record, table, document, and email referencing a data principal. The three execution paths — access, correction, erasure — all share this discovery layer.

The v2 design splits execution along two orthogonal axes: content type (structured versus Zone 1 unstructured) and retention status (expired versus active). The four resulting execution modes are: structured with expired retention (Delta DELETE plus VACUUM immediately); structured with active retention (block erasure, schedule and disclose the date); Zone 1 content (mask in place, de-index from search, schedule purge per collection retention); sensitive data with no retention obligation (erase immediately, crypto-shred backups).




FIGURE 8.1 Module 03 v2 with unified discovery combining Unity Catalog lineage and Zone 1 document search, three rights paths (access, correction, erasure), split execution by content type, combined erasure certificate, and residual retention register handoff to Module 06.

The residual retention register is the architectural expression of an important honesty commitment: the platform tells data principals exactly what was erased immediately, what was masked with scheduled final purge, and when each residual item will actually be deleted. This tripartite response — data export + erasure certificate + retention schedule — is what makes the hybrid execution defensible to the DPBI.

9 Module 04 — Breach detection and 72-hour alerting

The distinctive outcome is collapsing a 72-hour deadline into an 18-hour response, leaving 54 hours of buffer for legal review. Lakewatch's agentic detection plus Unity Catalog lineage plus Agent Bricks' notification drafting produces this margin.

The timeline in figure 9.1 tracks a worked scenario: an anomalous bulk export is detected by Lakewatch agents at T+0; the alert fires to the CISO team at T+3 minutes with a breach summary and the affected tables surfaced; Unity Catalog lineage resolves 82,341 affected principals by T+15 minutes; Agent Bricks auto-drafts the DPBI notification by T+2 hours; GC and CCO review during business hours; the notification is submitted at T+18 hours — 54 hours inside the Act's 72-hour deadline.

72-hour breach response timeline




FIGURE 9.1 Module 04 breach response timeline. Five stages from T+0 detection to T+18hr submission, leaving 75% of the 72-hour deadline as buffer.

The practical value of the buffer is not just compliance margin. It shifts legal review from a 2am emergency drafting exercise to a daylight-hours review of a populated template — which is how organizations survive breaches without burning out their GC team in the process.

10 Module 05 — Compliance audit and DPIA

The distinctive mechanism is that the DPIA doesn't need to be gathered because the metadata already exists in Unity Catalog. The scoring engine reads, weighs, and Agent Bricks composes — a three-month exercise becomes a three-day regeneration.

Module 05 is the reflective surface of the platform. It reads signal from Module 01 (the personal data register), Module 02 (consent coverage), Module 04 (breach readiness), and Unity Catalog metadata (processing activities), weights each against the twelve DPDP obligations with penalty-adjusted scoring, and produces four outputs: a live daily compliance score for the CCO's dashboard, a penalty-weighted gap heatmap for prioritization, the annual DPIA PDF in DPBI-ready format, and a remediation tracker with owners and SLAs.




FIGURE 10.1 Module 05 compliance scoring engine. Four inputs feed the engine; four outputs serve four audiences.

The scoring engine's weights are configurable in Unity Catalog, which matters because the DPDP Rules are still iterating — penalty structures will change, SDF designation criteria will firm up, and hardcoded weights would force re-architecture on every regulatory update. This design survives regulatory iteration explicitly.

11 Module 06 — Retention, masking and transfers

The distinctive mechanism is three parallel governance tracks — Delta DELETE plus crypto-shredding, Unity Catalog masking, and UC lineage-based cross-border monitoring — sharing one Workflows orchestration layer and four triggers.

The v2 design recognizes four trigger categories. Consent withdrawal fires when Module 02 records an opt-out. Retention expiry fires on scheduled policy boundaries. Employee exit fires 90 days after an HRIS status change. The fourth — DSR residual — is the most architecturally interesting: it's a deferred trigger set up when a DSR is resolved-with-residual and fires months or years later when the retention period elapses. Databricks Workflows' point-in-time scheduling against Delta metadata handles this naturally.




FIGURE 11.1 Module 06 v2 with four triggers (consent withdrawal, retention expiry, employee exit, DSR residual), three parallel execution tracks (physical deletion with crypto-shredding, UC column masking with tokenization vault, lineage-based cross-border monitoring), and the loop-close back to Modules 03 and 05.

Crypto-shredding in the left track handles the backup erasure problem that most DPDP platforms handle poorly or not at all. Active Delta data is cleaned by DELETE + VACUUM; records in offsite backups and archives are erased by destroying their per-record encryption keys — much cheaper than tape-sweep operations and fully defensible to regulators.

The loop-close box at the bottom is operationally important. When Module 06 eventually fires a scheduled DSR residual purge (years after the original request), it calls back to Module 03 to update the DSR record from "resolved with residual" to "fully closed," and updates Module 05's compliance score accordingly. The metric stops claiming partial resolution once the residual has cleared.

PART III

Addressing the realities

Most compliance platforms fail not on the technology but on the organizational prerequisites, the cross-functional friction points, and the shadow data surface that the tidy architecture diagrams quietly ignore. This part names those realities explicitly.

12 Module 0 — the organizational prerequisites

Four cross-cutting registries underpin every DPDP module. They are not platform features; they are organizational capabilities that must be built in parallel to the modules, not as a prerequisite that blocks them.

The **sources registry** is the master list of every system that holds personal data, with named owners. Module 01 reads from it; Module 03 uses it to locate data for DSRs; Module 06 uses it for cross-border monitoring. IT owns this registry with CCO validation.

The **legal hold registry** lists every principal whose data is under litigation, regulatory investigation, or tax audit. Modules 03 and 06 must check it before executing any erasure. Legal owns it; the registry's own data must itself be protected because the existence of a hold is often privileged.

The **vendor and processor registry** tracks every third party that handles personal data on the organization's behalf, with the DPA clauses that govern each relationship. Module 03 uses it to propagate erasure requests; Module 06 uses it for cross-border transfer monitoring. Procurement and GC co-own.

The **retention catalog** defines how long each data category must be kept, with legal basis. Module 06's triggers depend entirely on it. Recommend standardizing on 3-4 retention buckets (e.g., 3 years, 7 years, 10 years, indefinite-with-legal-basis) rather than bespoke per-field retention, which is a trap that stalls programmes for 6-12 months. Legal and Business co-own.




FIGURE 12.1 Module 0 prerequisites. Four registries each unlock multiple DPDP modules; Module 06 has the highest in-degree with four dependencies, which is why it is deployed in Phase 2 rather than Phase 1.

13 Friction audit across the six modules

Every module has non-technical friction points that determine whether the platform ships successfully. We audit them here by dominant angle — cost, legal, cross-team, regulatory, vendor — because the mitigation depends on which angle dominates.

Module 01

Shadow data is the bigger risk than the systems you know about (cross-team). Auto-discovery only works against systems IT is aware of. The long tail lives in department-managed spreadsheets, individual OneDrive folders, retired apps. Mitigation: pair Module 01 with a legal/compliance attestation survey; the automation surfaces candidates, humans complete the map.

Classification disputes with data owners (cross-team). Data owners contest whether a field is personal data; DPDP's definition is broad and tests are judgment calls. Mitigation: build a review UI where owners can dispute ai_classify's calls, with a compliance reviewer as tiebreaker.

LLM classification cost at scale (cost). Running ai_classify across billions of rows is expensive. Mitigation: sample-based classification with schema-level inference; reserve LLM functions for unstructured content and ambiguous structured fields.

Module 02

Pre-Act consent is the single biggest business risk (legal + business). Re-noticing millions of existing records risks mass opt-out. Mitigation: stage by customer value and engagement tier, lead with education messaging, use legitimate-interest grounds where the Act allows.

Touchpoint integration is program management, not engineering (cross-team). 10+ team engagements to update SDKs and UX. Mitigation: start with the two highest-volume channels covering 80% of events, prove the pattern, roll out over 6-12 months.

Granular consent destroys conversion rates (business). Strict per-purpose checkboxes tank signup conversion. Mitigation: tiered consent with core-service on legitimate-interest grounds, plus specific toggles only for purposes outside the core.

Module 03

Identity verification is a cost-per-request problem (cost + fraud risk). Mitigation: risk-tier the verification — email-link for access, step-up auth for erasure involving financial or biometric data.

Legal holds conflict with erasure mandates (legal). Mitigation: Legal Hold Registry checked pre-execution, implemented as Unity Catalog row-level tags.

Processor erasure propagation is contractual, not technical (vendor). Mitigation: standard DPA clauses requiring erasure verification, and the vendor registry tracks who holds what.

Module 04

"All breaches, regardless of severity" creates reporting noise (regulatory interpretation). Mitigation: agree a documented materiality threshold with Legal, aggregate below-threshold incidents in monthly summaries.

Cross-functional rehearsal is the gap most orgs fail on (cross-team). Mitigation: quarterly tabletop exercises using Lakewatch's simulation mode.

Lakewatch's Private Preview status creates architectural risk (vendor). Mitigation: pluggable detection engine — ship with MLflow Isolation Forest, swap in Lakewatch when available.

Module 05

DPDP Rules are still being notified (regulatory). Build the scoring engine with configurable weights in Unity Catalog; design for iteration.

SDF designation uncertainty (regulatory). Architect for SDF obligations from the start; retrofitting is more expensive than over-engineering.

Board dashboards compete for executive attention (business). Tie compliance metrics to business metrics the board already tracks.

Module 06

Defining retention periods is a 6-12 month project (cross-team). Mitigation: standardize on 3-4 retention buckets; force categories into them. Bespoke per-field retention is a trap.

Masking breaks legitimate analytics (business). Mitigation: tokenization with controlled re-identification via purpose-coded audited unmasking.

Backup and archive erasure is the dirty secret of DPDP compliance (legal + cost). Mitigation: crypto-shredding — encrypt per-record, erase the key.

14 Shadow data and the three zones

Most PII that causes real breaches lives in collaboration tools, emails, and personal devices — not in the structured databases Module 01 naturally targets. The response splits into three zones by scannability, and only one of them is an extension of Module 01.




FIGURE 14.1 Three zones of shadow data. Zone 1 is in scope for direct scanning; Zone 2 requires security tooling (CASB/DLP/MDM); Zone 3 is off-limits and covered by policy and attestation only.

Zone 1 — sanctioned enterprise cloud. SharePoint, OneDrive, Exchange Online, Google Workspace, Teams and Slack Enterprise archives. The company legally owns this data; Module 01 extends here via Lakeflow Connect's SharePoint connector, Microsoft Graph API, and Google Workspace APIs. The key governance detail is that scanning employee mailboxes for customer PII is itself a processing activity requiring a lawful basis (typically legitimate purpose, documented). Exclusions for privileged content, HR-sensitive material, and legal holds must be explicit and GC-signed before scanning starts.

Zone 2 — the edge. Corporate laptops, managed mobile devices, shadow SaaS apps, outbound email with PII attachments. Here the goal shifts from discovery to prevention. The tooling is different — CASB for shadow SaaS, DLP for exfiltration blocking, MDM for endpoints — and the owner is the CISO, not the data team. The right relationship with Module 01 is that Zone 2 findings feed the register: when CASB discovers an unsanctioned tool, it gets onboarded to Zone 1 or blocked.

Zone 3 — off-limits. Personal email accounts, personal cloud storage, BYOD personal partitions, paper documents. Scanning is either legally impossible, practically impossible, or privacy-invasive enough that the cost-benefit is upside-down. Controls here are policy, attestation, and training — HR and Legal own them.

The practical implication: the proposal should scope Module 01 as Zone 1 expansion explicitly and flag Zones 2 and 3 as parallel workstreams requiring their own sponsors. If the proposal tries to bundle all three under one budget line, it stalls on "whose budget is this?" conversations.

PART IV

Operational walkthroughs

Six workflows together cover the complete operating envelope of a DPDP programme — what happens when things go wrong, how the programme evidences itself, how individual principals are served, how the regulator is answered, how the consent foundation is built, and how it is sustained day-to-day. Each touches multiple modules; no module can serve any workflow alone.

15 Breach response timeline

The workflow with the tightest clock. Five stages compressing a 72-hour deadline into an 18-hour response, reviewed in detail under Module 04 above. See figure 9.1 for the timeline and architectural detail.

The operational point worth reiterating here is the CISO-to-GC handoff at T+2 hours. The CISO owns detection and scoping; the GC owns the notification draft. The diagram shows both as primary owners in consecutive phases — the visual reminder that ownership transfers mid-workflow, and that this specific moment requires rehearsal. Organizations that discover the handoff for the first time during a real incident spend hours getting Legal caught up on technical detail while the clock runs. Quarterly tabletop exercises fix this.

16 Compliance audit cycle

The workflow with the longest cadence — daily scoring, quarterly board packs, annual DPIA. Unlike breach which is reactive and concentrated, this is proactive and rhythm-based.

The CCO owns the workflow throughout. What makes this cycle distinctive is that the annual DPIA phase has three primary owners in succession: the CCO owns the DPIA as a compliance artifact; the CFO owns the ROI narrative that the board cares about; the GC owns the legal review because it is a regulator filing. Each owns a different section of the same document. The review cadence must accommodate all three signatures, which is why the cycle runs over weeks rather than a single sign-off moment.

The automation in Agent Bricks compresses the DPIA drafting phase dramatically — three months of manual compilation becomes three days of agent-composed drafts against Unity Catalog metadata. But the review phase cannot be compressed the same way; legal and business judgment over lawful basis and processing necessity is where human review is load-bearing.

17 Single DSR — Priya's walkthrough

A realistic scenario — an ex-employee 18 months post-departure requesting access plus erasure — touches every zone (structured HR data, Zone 1 emails and documents, biometric attendance), crosses every module, and surfaces every policy decision.

Priya left the company 18 months ago. She was a senior account manager for 6 years. She submits a DPDP request today asking for both access to all her personal data and erasure of everything not legally required to be retained. No litigation, no investigation — a legitimate rights exercise.




FIGURE 17.1 Priya's end-to-end DSR walkthrough. Discovery returns 47 assets across 9 systems; execution splits into four retention-aware modes; the 5-day response bundle includes data export, erasure certificate, and a disclosed retention schedule for residuals that will fire through Module 06 in 2032.

Three things the walkthrough reveals that the module diagrams alone don't. First, **execution splits along two orthogonal axes** — content type (structured versus Zone 1) and retention status (expired versus active) —

producing four distinct execution modes. Tax records under Income Tax Act mandate block immediate erasure; performance reviews have no retention obligation and are deleted immediately; biometric data gets crypto-shredded through to backups; Zone 1 emails are masked and de-indexed with scheduled physical purge. Second, **the response bundle is tripartite, not unitary**: data export, erasure certificate, retention schedule. The schedule disclosure is the honesty commitment that makes the hybrid execution defensible. Third, **the workflow extends years beyond the response date**; the scheduled residual purges fire through Module 06 in December 2032 — 6.5 years later — with automated final certificates delivered at that point.

18 DPBI Section 28 inspection

The most high-stakes external-facing workflow. A regulator notice asking for evidence of 2026-2027 DSR handling must be answered with a reproducible packet inside 30 days. The platform's real job is reproducibility, not dashboards.

The scenario: Q2 2028, the DPBI issues a Section 28 notice requesting a random sample of 50 DSRs plus 6 complaint cases from calendar years 2026 and 2027. Non-compliance or materially incomplete submissions can trigger penalties up to ₹50 crore — not for the original handling, which may have been fine, but for failure to produce evidence.




FIGURE 18.1 DPBI Section 28 inspection response. Cross-module evidence query reads from all six modules; Delta time travel reproduces historical state; Agent Bricks composes the 56-case packet; GC, CCO, CEO review and sign-off within 16 days of notice.

Delta time travel is the load-bearing component. Every table's state can be reproduced as of any requested date within the configured retention window, which is how the evidence query answers "what did you know on 12 May 2026 when you handled this DSR?" The Unity Catalog audit log records who queried what, when, with which purpose code. Together these produce regulator-grade reproducibility — the DPBI can ask any follow-up question and the platform answers from the same data that served the original handling.

The GC has primary ownership across all four phases of this workflow — unique to this inspection response. That is because a Section 28 response is a legal proceeding, not an operational workflow. Every decision carries legal implications that the GC must own.

19 Retrospective consent campaign

The programme's founding event. 8.4 million existing records re-noticed over 4 months. A one-shot exercise with multi-year consequences — the consent baseline every other workflow depends on.




FIGURE 19.1 Retrospective consent campaign across 4 months. Segmentation into 12 cohorts, 22-language templates with legal approval, four waves (pilot, active, dormant, no-consent), three response categories, and steady state as the CMO's consented-audience foundation.

Four design decisions in this workflow are worth naming. First, **the pilot wave is non-negotiable** — sending to 50,000 most-engaged customers before the 3.8M active wave validates response-rate assumptions on a recoverable scale. Second, **Wave 4 is categorically different** from Waves 1-3: it's 2.2M records without any consent trail, often from partner or lead-gen origins, where sending a re-notice is itself DPDP-questionable; expect 60-80% to purge rather than convert, which is a realistic outcome, not a failure mode. Third, **"default to opt-out after 60 days"** is a documented legal-interpretation decision, not a platform default — GC must sign it off in the Module 0 interpretation framework. Fourth, **Module 02's consented-audience API stays live throughout**, so marketing shrinks during the transition rather than freezing — which is what prevents revenue collapse during the 4-month window.

The CHO inherits this pattern for the employee population. The playbook built for customers — segmentation, tiered messaging, waves, non-response handling — maps directly onto re-noticing the existing workforce for employment-data consent. Running the customer campaign well means having the playbook ready for the employee one.

20 New-principal consent — Meera's walkthrough

The steady-state intake funnel that runs every day forever. Unlike the one-shot retrospective campaign, this is the continuous flow that either builds compliant trust or erodes it.

Meera, 28, opens a savings account via the mobile app — first interaction with the company. Her walkthrough in figure 20.1 exercises the full new-principal pipeline: age gate first, notice rendered in Kannada at version 4.2, tiered consent capture with core service on legitimate-interest grounds and three toggles for optional purposes, atomic Lakebase write committing profile plus three consent events as one transaction, propagation within five minutes to the Delta log, Module 01 register, and Module 06 retention binding.




FIGURE 20.1 Meera's new-principal onboarding. Age gate → notice → tiered consent capture → atomic Lakebase write → propagation. Each toggle is a discrete consent event; the retention clock starts at the consent moment per purpose.

The load-bearing invariant is **channel parity**: web, branch kiosk, call center IVR, partner referral — every channel varies only in UX; all write to the same Lakebase endpoint with the same schema. Without this invariant enforced at the CTO level, the platform silently accumulates six different consent event schemas with four different quality levels, and three years later the register has drifted past recoverability. With it enforced, every record in the system traces back to a specific notice version rendered in a specific language at a specific moment — or it doesn't, in which case the record predates the platform and is handled by the retrospective campaign. There are no orphans; every principal has provenance.

The age gate is the single most important control in the flow. Under-18 declaration routes to a parental consent subsystem that is its own programme of work if the business serves children; for most enterprises the safest stance is "we don't knowingly onboard under-18 principals" with the age gate as a hard block. Product-line-level decision, not a platform-level one.

PART V

The coverage story

Two artifacts close the structural argument. Persona journeys across the six workflows show who engages when; the workflow-to-module coverage matrix shows what the platform does and in what order the modules should be deployed.

21 Persona journeys across the six workflows

Six mini-matrices, one per workflow. Each shows all seven personas against the workflow's four phases, distinguishing primary owners from active participants from the informed-only. Together they reveal the multi-threading pattern.




FIGURE 21.1 Breach response persona journey. CISO owns throughout; GC takes primary ownership at draft phase; CMO and CHO gated by affected-population criterion.




FIGURE 21.2 Compliance audit cycle persona journey. CCO owns the cycle; three primary owners converge at the annual DPIA (CCO, CFO, GC); domain personas rotate in quarterly to review their section.




FIGURE 21.3 Single DSR persona journey. Four primary-owner cells — CCO, CHO, GC, CTO — indicate a workflow that fundamentally needs cross-functional execution.

Workflow · DPBI Section 28 inspection response




FIGURE 21.4 DPBI inspection persona journey. GC has primary ownership across all four phases (unique to this workflow); CFO is active in every phase due to penalty exposure.

Workflow · retrospective consent campaign (4 months)




FIGURE 21.5 Retrospective consent persona journey. CMO owns all four phases; CHO inherits the pattern at the steady-state phase for the employee campaign.

Workflow · new principal consent (continuous)

	Framework	Capture UX	Operation	Monitor
GC	Framework	Approves UX	Escalations	Version rev.
CTO	Schema	SDK review	Operates	Drift alerts
CMO	Channels	Owns UX	Conversion	Optimises
CCO	Aligns	Reviews	Coverage	Reports
CHO	Observes	Observes	Employee ops	HR coverage
CISO	—	Bot / fraud	Monitors	Informed
CFO	—	—	Cost / capt.	Unit econ.

■ Primary owner
 ■ Active participant
 ■ Informed / consulted

FIGURE 21.6 New-principal consent persona journey. The most distributed ownership of any workflow — GC, CTO, CMO, CCO, and CHO all have at least one primary-owner cell.

Three cross-workflow patterns matter for the multi-threading strategy. The **CCO is primary or active in every single workflow** — across 24 cells, never marked with a dash; this is the operational truth behind the CCO being the platform's core persona. The **GC and CFO appear as active or primary in every workflow** too, with different intensity patterns. These three — CCO, GC, CFO — are the "always in the room" personas the proposal should lead with. The **CISO, CMO, CHO, CTO rotate in and out**, each having at least one workflow where they are primary but several where they are uninvolved. This is how to pace the multi-threaded outreach: lead with CCO-GC-CFO for breadth, bring in domain personas for their specific workflows.

22 Workflow-to-module coverage matrix

The structural dependency story. Six workflows × six modules, grading each cell as Core (workflow breaks without module), Support (meaningful participation), or Touch (queried or updated). Read-versus-write direction flags data-flow dependencies.




FIGURE 22.1 Workflow-to-module coverage. Module 06 holds the most core cells (3); Module 05 has the highest total hit count (6) as the reflective surface. The bottom-right cluster of M05 + M06 across new-principal and retrospective consent workflows is where sustained operational value lives.

Five observations the matrix surfaces that matter for proposal sequencing.

Module 06 is the busiest module — 3 core cells and 5 total hits — because it is core to erasure decisions (single DSR), opt-out purges (retrospective consent), and retention binding at capture (new principal). The pitch guide classifies Module 06 under CISO and GC ownership, which is correct for governance but understates operational centrality.

Module 05 has the highest total hit count — 6 — but only 1 core cell. Module 05 reads from every other module and writes the aggregate score; the compliance audit cycle is its home workflow, everywhere else it is supporting. For the CFO conversation this is concrete: every workflow feeds the score.

Module 04 is the narrowest — 1 core cell, 3 total hits. Breach response is its only core dependency. This means Module 04 can be built or deferred largely independently of the other modules; if Lakewatch's Private Preview timeline creates architectural risk, the rest of the platform is insulated.

Modules 01 and 02 show up in 5 workflows each, tied for second-busiest, validating the "lead with 01 and 02" recommendation empirically. Every workflow except purely-internal breach response needs either the data register or the consent store.

The breach workflow is deliberately sparse (3 hits) — this is a design feature, not a gap. Every cross-module dependency is a potential latency source at T+0; the minimal set is the right design.

The sequencing implication: **Phase 1** stands up Modules 01 and 02, unblocking 5 of the 6 workflows at minimum viable coverage. **Phase 2** brings Module 06 online — non-intuitive relative to the pitch guide's ordering but correct given the matrix, because Module 06 is the most-cited core module and without it Phase 1's consent capture and data register become write-only. **Phase 3** stands up Module 05 for aggregate scoring. **Phase 4** brings Module 03 to full scope with Zone 1 rights handling. **Phase 5** completes Module 04 with Lakewatch, following the Private Preview release cycle. Every phase ships a working platform, never half-working.

PART VI

The proposal

The 14-day rapid discovery sprint is the entry ticket; Phase 1 is what it unlocks. This part presents the POC scope card, its archetype-specific shading, the Databricks internal classification strategy, and the Phase 1 outline.

23 POC scope card — 14-day sprint

One source system. Two modules. Three personas demoed to in the same room. The sprint proves PII discovery, consent capture, and synthetic DSR end-to-end against live data.

POC scope · 14-day rapid discovery sprint

Objective

Prove PII discovery · consent capture · DSR fulfilment
against a live source system in 14 days

Scope

One source
CRM or HRMS
Structured only
One-time export

Two modules
Module 01 · PII
Module 02 · consent
Minimum viable

Three personas
CCO · register
CMO · consent
GC · DSR proof

Deliverables

Data register
All PII columns
classified, tagged
and queryable

Consent flow
1000 synthetic
events captured
withdrawal under 5 min

Synthetic DSR
Access + erasure
full bundle
and certificate

Success criteria

Register: every table scanned · precision $\geq$ 90% on spot-check · lineage present

Consent: 1000 events in Lakebase · Delta log immutable · withdrawal under 5 min propagation

DSR: synthetic principal discovered · erasure executed + verified · certificate produced

Demo: all three artifacts visible on Day 14 to CCO, CMO, GC, CFO in one session

Team

Databricks side

1 solution architect
2 field engineers

Customer side

1 source owner (10hr/wk)
1 compliance SME · GC liaison

Timeline

Day 1-3

Ingest + classify

Day 4-7

Consent store

Day 8-11

DSR end-to-end

Day 12-14

Demo + decision

Exit decisions at Day 14

Go · Phase 1 approved

Scope + budget sign-off

Iterate · extended POC

Adjust and re-demo in 2 weeks

Explicit exclusions

Out of POC · Zone 1 ingestion · breach detection · annual DPIA
Cross-border monitoring · employee DSR · retention catalog
Legal hold registry · vendor processor registry
All belong to Phase 1 · not the sprint

FIGURE 23.1 POC scope card. Objective, scope boundaries, three deliverables mapped to three personas, success criteria, team composition, four-phase timeline, exit decisions, and explicit exclusions.

Five design decisions in the card worth naming.

One source system is non-negotiable. Multiple sources would inherit the 40+-team credential-coordination bottleneck on Day 1 and fail on organizational rather than technical grounds. One source with a source owner at 10 hours per week is the manageable size.

Structured only, no Zone 1. Zone 1 needs GC sign-off on the exclusion policy — realistically 4-6 weeks. A POC that includes Zone 1 either rushes the policy review or bursts the 14-day window. Zone 1 belongs to Phase 1 proper with its own kickoff.

Modules 01 and 02 only, not 05 or 06. Module 05's score is only meaningful when aggregating multiple working modules; running it against a single POC source produces a placeholder number that erodes trust. Module 06 depends on the retention catalog which hasn't been built. Three modules proven is stronger than six half-proven.

The synthetic DSR is the single most important deliverable. It is the integration test proving Modules 01 and 02 work together rather than merely existing in parallel. Two separate artifacts read as two tools; one end-to-end DSR reads as one platform.

The demo on Day 14 requires all four personas in the same room. Separate per-persona demos let each see only their slice and undermine the one-platform story. The 45-minute joint session is itself a test of organizational commitment.

The exclusions list at the bottom of the card is the most important section. Every excluded item represents a topic a stakeholder will raise during kickoff with an implicit expectation that it is included. The CISO will ask about breach detection; the CHO about employee DSR; the CFO about the full compliance score. Naming exclusions explicitly prevents scope creep through social pressure.

24 POC archetype comparison

Three realistic customer archetypes shade the POC shape differently. BFSI with existing UC; telco without UC; retail AWS-native new-logo. Each changes the classification, product attach weights, and Phase 1 scope.

POC shape by customer archetype

	BFSI	Telco	Retail
Primary classification	UC extension Existing UC	UC adoption DPDP driver	New-logo Full adoption
Agent Bricks weight	Heavy SDF DPIA	Medium DPIA plus lang	Medium Notice gen
Lakebase weight	Heavy Consent OLTP	Architectural High volume	Medium Consent + DSR
Lakeflow weight	Light Existing pipes	Heavy Volume	Medium SaaS heavy
Source system	Core banking or policy admin	Subscriber mgmt Single CRM	E-comm CRM or main DB
Stakeholder complexity	High Legal 15 hr/wk	CRO + CCO split roles	Board / PE driven
Phase 1 scope	DPDP-focused Extension	DPDP + UC expansion	Full adoption DPDP first

Heaviest consideration
 Standard consideration
 Light touch

FIGURE 24.1 POC archetype comparison. Six dimensions shaded per archetype; coral cells mark heaviest considerations.

BFSI with existing Unity Catalog. Primary classification shifts to "UC extension" rather than UC governance because the customer already runs UC. Agent Bricks becomes the heaviest attachment (SDF-mandated DPIA automation); Lakebase is heavy (consent OLTP they likely don't have); Lakeflow is light (existing ingestion patterns). Source system is core banking or policy admin, not generic CRM. Legal liaison engagement 15 hours per week. Add exclusion for RBI Account Aggregator obligations. Phase 1 is DPDP-focused extension.

Large telco without existing UC. Primary classification is "UC adoption driven by DPDP" — actually the cleanest classification because it is a UC adoption opportunity. Lakeflow becomes heavy (telco volumes non-trivial); Lakebase becomes architectural decision (subscriber-scale consent volumes); stakeholder complexity flagged as CRO/CCO split (regulatory affairs separate from compliance). Source system is subscriber base management. Scale consent-event target from 1000 to 10,000.

Retail or consumer goods, AWS-native, new-logo. Primary classification is new-logo Databricks adoption with DPDP as the entry use case. Day 1-3 expands to include workspace provisioning and UC setup; add a platform architect to the Databricks team; add a customer cloud engineer for IAM and networking decisions. Phase 1 is the most consequential of the three — full Databricks adoption with DPDP as the first workload,


adjacent analytics workloads flagged as expansion pipeline. Extra care during POC to prevent Phase 1 scope expansion; let the Phase 1 conversation shape adjacent opportunities separately.

25 Databricks use case classification


Internal to Databricks, the POC should classify across three tiers rather than a single label, with an emerging-category argument for Regulatory Compliance Platform as a repeatable vertical motion.

Use case classification

Primary



Secondary · product attach



Tertiary · expansion track



Emerging · new category proposal




FIGURE 25.1 Three-tier classification for the DPDP POC within Databricks use case taxonomy. Primary, secondary product attach, tertiary expansion, and an emerging-category proposal.

Primary — Unity Catalog data governance. The classification that aligns the POC to existing Databricks internal motions and field staffing. UC-led opportunities are well-understood; existing playbooks, deployment patterns, and success metrics apply. Every other compliance capability in the platform is a consumer of UC's tags, lineage, and audit log — structurally honest framing.

Secondary — Agent Bricks, Lakebase, Lakeflow Connect. Product-attach tags that increase internal opportunity profile, unlock product-line specialists, and flow revenue attribution correctly. Each product has its own field motion that co-sells naturally with a governance-led opportunity.

Tertiary — Lakewatch and Genie. Expansion opportunities that don't feature in the 14-day POC itself but are natural Phase 2 extensions. Keeping them tertiary preserves forward visibility without overloading the POC commitment.

Emerging — Regulatory Compliance Platform. DPDP in India follows a pattern established by GDPR and being replicated across US state laws, Brazil's LGPD, Saudi Arabia's PDPL, Singapore's PDPA. The platform architecture is not DPDP-specific; it is a regulatory compliance pattern with DPDP as the current instance. Proposing it as an emerging category positions the customer as a lead reference for a broader pattern.

A specific recommendation on framing: *internally to Databricks*, use the three-tier classification. *To the customer*, this is a DPDP compliance platform programme — full stop. The Databricks product attachments are implementation details that feature in technical sections, not in the compliance outcome framing. Keep the two separate.

26 Phase 1 outline and next steps

What comes out of the Day 14 sign-off. The 3-6 month commitment that follows the POC, phased across five deployment waves, with Module 0 as a parallel organizational workstream.

Phase 1 (months 1-3) — Modules 01 and 02 stood up at full scope. Structured ingestion across the prioritized source list; consent capture machinery integrated with the two highest-volume channels (typically web plus mobile); Lakebase, Delta log, and DLT propagation live. The CCO's register and the CMO's consented-audience API become real operational assets.

Phase 2 (months 3-6) — Module 06 online. Retention catalog signed off by Legal in parallel; workflow triggers live; crypto-shredding configured for backups; UC column masking for governed analytics. DSR erasure and retention binding for new principals become executable end-to-end.

Phase 3 (months 6-9) — Module 05 aggregating. Compliance score live on CCO dashboard; Agent Bricks DPIA drafting configured; board pack template instantiated. Score becomes a live metric.

Phase 4 (months 9-12) — Module 03 full scope including Zone 1. SharePoint and Exchange ingestion live under the GC-signed exclusion policy; split execution for Zone 1 erasure operational; residual retention register integrated with Module 06.

Phase 5 (months 12-18) — Module 04 with Lakewatch. Dependent on Lakewatch Private Preview exit; interim MLflow Isolation Forest fallback available if Lakewatch timeline slips; pluggable detection engine design means the downstream pipeline is unaffected either way.

Parallel — Module 0 organizational workstream. Source registry, legal hold registry, vendor and processor registry, retention catalog. Four different owners outside the platform team; each needs its own sponsor and kickoff on Day 1. The DPDP Legal Interpretation Framework (materiality threshold, non-response policy, legitimate-interest grounds, erasure-versus-retention decisions) bundled as a single GC-owned document signed off before the affected modules ship.

Immediate next steps before the Day 1 of the POC. Source system selected and signed off; source owner named with confirmed 10-hour weekly capacity; compliance SME and GC liaison confirmed; one-time CSV export scheduled with the source team; Databricks workspace provisioned with Unity Catalog, Lakebase, and Agent Bricks enabled; exploratory data use agreement signed covering the 14-day window; joint CCO/CMO/GC/CFO Day-14 demo slot placed on calendars.

Most POCs slip on these readiness items rather than execution. Naming them explicitly lets the account team drive them to closure in the 2-3 weeks before Day 1.

End of document · 26 sections, 24 figures, six walkthroughs, six modules, one platform.